Build Smarter AI Apps: Empower LLMs with LangChain

Module Cheat Sheet: Introduction to LangChain in GenAI

Package/Method	Description	Code Example

WatsonxLLM	A class from the ibm_watson_machine_learning.foundation_models.extensions.langchain module that creates a LangChain compatible wrapper around IBM's watsonx.ai models.	```from ibm_watsonx_ai.foundation_models import ModelInference```
		 ```from ibm_watson_machine_learning.foundation_models.extensions.langchain import WatsonxLLM```
		
		 ```model_id = 'mistralai/mixtral-8x7b-instruct-v01'```
		 ```parameters = {```
		 ```GenParams.MAX_NEW_TOKENS: 256,```
		 ```GenParams.TEMPERATURE: 0.2,```
		 ```}```
		 ```credentials = {"url": "https://us-south.ml.cloud.ibm.com"}```
		 ```project_id = "skills-network"```

		``python
model = ModelInference(
    model_id=model_id,
    params=parameters,
    credentials=credentials,
    project_id=project_id
)

mixtral_llm =
WatsonxLLM(model=model)
response =
mixtral_llm.invoke("Who is
man's best friend?")
```
Copied!Wrap Toggled! |
| **Message Types** | Different types of messages that chat models can use to provide context and control the conversation. The most common message types are SystemMessage, HumanMessage, and AIMessage. | ```python
from
langchain_core.messages
import HumanMessage,
SystemMessage, AIMessage


msg = mixtral_llm.invoke([

SystemMessage(content="You are a helpful AI bot that assists a user in choosing the perfect book to read in one short sentence"),

HumanMessage(content="I enjoy mystery novels, what should I read?")
])
```
Copied!Wrap Toggled! |

| **PromptTemplate** | A class from the langchain_core.prompts module that helps format prompts with variables. These templates allow you to define a consistent format while leaving placeholders for variables that change with each use case. | ```python
from langchain_core.prompts import PromptTemplate


prompt = PromptTemplate.from_template("Tell me one {adjective} joke about {topic}")
input_ = {"adjective": "funny", "topic": "cats"}

formatted_prompt = prompt.invoke(input_)
```<br>Copied!Wrap Toggled! |
| **ChatPromptTemplate** | A class from the langchain_core.prompts module that formats a list of chat messages with variables. These templates consist of a list of message templates themselves. | ```python
from langchain_core.prompts import ChatPromptTemplate


prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful assistant"),
    ("user", "Tell me a joke about {topic}")
])

input_ = {"topic": "cats"}
formatted_messages = prompt.invoke(input_)
```<br>Copied!Wrap Toggled! |

MessagesPlaceholder	A placeholder that allows you to add a list of messages to a specific spot in a ChatPromptTemplate. This capability is useful when you want the user to pass in a list of messages you would slot into a particular spot.	```from langchain_core.prompts import MessagesPlaceholder from langchain_core.messages import HumanMessage prompt = ChatPromptTemplate.from_messages([("system", "You are a helpful assistant"), MessagesPlaceholder("msgs")]) input_ = {"msgs": [HumanMessage(content="What is the day after Tuesday?")]} formatted_messages = prompt.invoke(input_)``` Copied!Wrap Toggled!
JsonOutputParser	A parser that allows users to specify an arbitrary JSON schema and query LLMs for outputs that conform to that schema. A parser is useful for obtaining structured data from LLMs.	```from langchain_core.output_parsers import JsonOutputParser from langchain_core.pydantic_v1 import BaseModel, Field class Joke(BaseModel): setup: str =```

		Field(description="question to set up a joke")
		    punchline: str = Field(description="answer to resolve the joke")
		
		output_parser = JsonOutputParser(pydantic_object=Joke)
		
		format_instructions = output_parser.get_format_instructions()
		prompt = PromptTemplate(
		    template="Answer the user query.\n{format_instructions}\n{query}\n",
		    input_variables=["query"],
		
		partial_variables={"format_instructions": format_instructions},
		)
		
		chain = prompt | mixtral_llm | output_parser
		
		Copied!Wrap Toggled!

CommaSeparatedListOutputParser	A parser used to return a list of comma-separated items. This parser converts the LLM's response into a Python list.	from langchain.output_parsers import CommaSeparatedListOutputParser output_parser = CommaSeparatedListOutputParser()

		``
format_instructions =
output_parser.get_format_in
structions()
prompt = PromptTemplate(
    template="Answer the user
query.
{format_instructions}\nList
five {subject}.",

    input_variables=["subject"],

    partial_variables={"format_in
structions":
format_instructions},
)

chain = prompt | mixtral_llm |
output_parser
result =
chain.invoke({"subject": "ice
cream flavors"})
```

Copied!Wrap Toggled! |
| **Document** | A class from the langchain_core.documents module that contains information about some data. This class has the following two attributes: page_content (the content of the document) and metadata (arbitrary metadata associated with the document). | ```
from
langchain_core.documents
import Document


doc = Document(
    page_content="""Python is
an interpreted high-level
general-purpose
programming language.
        Python's design
philosophy emphasizes code
readability with its notable
use of significant
``` |

		indentation."""`,
		  metadata={
		    'my_document_id' :
		234234,
		    'my_document_source' :
		"About Python",
		
		'my_document_create_time'
		: 1680013019
		   }
		)
		
		Copied!Wrap Toggled!

PyPDFLoader	A document loader from the langchain_community.document_loaders that loads PDFs into Document objects. You can use this document loader to extract text content from PDF files.	```from langchain_community.document_loaders import PyPDFLoader```
		
		```loader = PyPDFLoader("path/to/document.pdf")```
		```documents = loader.load()```
		
		Copied!Wrap Toggled!

WebBaseLoader	A document loader from the langchain_community.document_loaders that loads content from websites into Document objects. You can use this document loader to extract text content from web pages.	```from langchain_community.document_loaders import WebBaseLoader```
		
		```loader = WebBaseLoader("https://python.langchain.com/v0.2/docs/introduction/")```
		```web_data = loader.load()```

CharacterText Splitter	A text splitter from langchain.text_splitter that splits text into chunks based on characters. This splitter is useful for breaking long documents into smaller, more manageable chunks for processing with LLMs.	```python
		from langchain.text_splitter import CharacterTextSplitter
		
		text_splitter = CharacterTextSplitter(
		    chunk_size=200, # Maximum size of each chunk
		    chunk_overlap=20, # Number of characters to overlap between chunks
		    separator="\n" # Character to split on
		)
		chunks = text_splitter.split_documents(documents)
		``

RecursiveCharacterTextSplitter	A text splitter from langchain.text_splitter that splits text recursively based on a list of separators. This splitter tries to split on the first separator, then the second separator, and any subsequent separators, until the chunks of text attain the specified size.	```python
		from langchain.text_splitter import RecursiveCharacterTextSplitter
		
		text_splitter = RecursiveCharacterTextSplitter(
		    chunk_size=500,
		    chunk_overlap=50,
		    separators=["\n\n", "\n", ". ", " ", ""]
		)
		chunks =
		``

		text_splitter.split_document s(documents) Copied!Wrap Toggled!

WatsonxEmbe ddings	A class from langchain_ibm that creates embeddings (vector representations) of text using IBM's watsonx.ai embedding models. You can use these embeddings for semantic search and other vector-based operations.	from langchain_ibm import WatsonxEmbeddings
		
		from ibm_watsonx_ai.metanames import EmbedTextParamsMetaNam es
		
		embed_params = {
		
		EmbedTextParamsMetaNam es.TRUNCATE_INPUT_TOKE NS: 3,
		
		EmbedTextParamsMetaNam es.RETURN_OPTIONS: {"input_text": True},
		}
		
		watsonx_embedding = WatsonxEmbeddings(
		  model_id="ibm/slate-125m-english-rtrvr",
		  url="https://us-south.ml.cloud.ibm.com",
		  project_id="skills-network",
		  params=embed_params,
		)
		
		Copied!Wrap Toggled!

Chroma	A vector store from langchain.vectorstores that stores embeddings and provides methods for similarity search. You can use Chroma for storing and retrieving documents based on semantic similarity.	```from langchain.vectorstores import Chroma```
		
		 ```// Create a vector store from documents```
		 ```docsearch = Chroma.from_documents(chunks, watsonx_embedding)```
		
		 ```// Perform a similarity search```
		 ```query = "Langchain"```
		 ```docs = docsearch.similarity_search(query)```
		
		 Copied!Wrap Toggled!

Retrievers	Interfaces that return documents given an unstructured query. Retrievers accept a string query as input and return a list of Document objects as output. You can use vector stores as the backbone of a retriever.	```# Convert a vector store to a retriever```
		
		 ```retriever = docsearch.as_retriever()```
		
		 ```// Retrieve documents```
		 ```docs = retriever.invoke("Langchain")```
		
		 Copied!Wrap Toggled!

ParentDocumentRetriever	A retriever from langchain.retrievers that splits documents into small chunks for embedding but returns the parent documents during retrieval. This retriever balances accurate	```from langchain.retrievers import ParentDocumentRetriever```
		
		 ```from langchain.storage import InMemoryStore```

| | embeddings with context preservation. | ```python
parent_splitter = CharacterTextSplitter(chunk_size=2000, chunk_overlap=20)
child_splitter = CharacterTextSplitter(chunk_size=400, chunk_overlap=20)

vectorstore = Chroma(

    collection_name="split_parents",

    embedding_function=watsonx_embedding
)

store = InMemoryStore()

retriever = ParentDocumentRetriever(
    vectorstore=vectorstore,
    docstore=store,

    child_splitter=child_splitter,

    parent_splitter=parent_splitter,
)

retriever.add_documents(documents)
retrieved_docs = retriever.invoke("Langchain")
```
Copied!Wrap Toggled! |
|---|---|---|
| **RetrievalQA** | A chain from langchain.chains that answers questions based on retrieved documents. The | ```python
from langchain.chains import RetrievalQA
``` |

	RetrievalQA chain combines a retriever with an LLM to generate answers based on the retrieved context.	```
		qa = RetrievalQA.from_chain_type(
		    llm=mixtral_llm,
		    chain_type="stuff",
		
		retriever=docsearch.as_retriever(),
		
		return_source_documents=False
		)
		query = "what is this paper discussing?"
		answer = qa.invoke(query)
		`` Copied!Wrap Toggled!

ChatMessage History	A lightweight wrapper from langchain.memory that provides convenient methods for saving HumanMessages, AIMessages, and then fetching them all. You can use the ChatMessageHistory wrapper to maintain conversation history.	```
		from langchain.memory import ChatMessageHistory
		
		history = ChatMessageHistory()
		history.add_ai_message("hi!")
		history.add_user_message("what is the capital of France?")
		// Access the messages
		history.messages
		// Generate a response using the history
		ai_response =
		``

		mixtral_llm.invoke(history.messages) Copied!Wrap Toggled!

ConversationBufferMemory	A memory module from langchain.memory that allows for the storage of messages and conversation history. You can use this memory module conversation chains to maintain context across multiple interactions.	from langchain.memory import ConversationBufferMemory
		
		from langchain.chains import ConversationChain
		
		conversation = ConversationChain(
		   llm=mixtral_llm,
		   verbose=True,
		
		memory=ConversationBufferMemory()
		)
		response = conversation.invoke(input="Hello, I am a little cat. Who are you?")
		
		Copied!Wrap Toggled!

LLMChain	A basic chain from langchain.chains that combines a prompt template with an LLM. It's the simplest form of chain in LangChain.	from langchain.chains import LLMChain template = """Your job is to come up with a classic dish from the area that the users suggests. {location}

		YOUR RESPONSE:

		""
		
		 prompt_template = PromptTemplate(template=template, input_variables=['location'])
		
		 location_chain = LLMChain(
		   llm=mixtral_llm,
		   prompt=prompt_template,
		   output_key='meal'
		)
		
		 result = location_chain.invoke(input={'location':'China'})
		
		 Copied!Wrap Toggled!

SequentialChain	A chain from langchain.chains that combines multiple chains in sequence, where the output of one chain becomes the input for the next chain. SequentialChain is useful for multi-step processing.	from langchain.chains import SequentialChain
		
		 // First chain - gets a meal based on location
		 location_chain = LLMChain(
		   llm=mixtral_llm,
		
		 prompt=location_prompt_template,
		   output_key='meal'
		)
		 // Second chain - gets a recipe based on meal
		 dish_chain = LLMChain(
		   llm=mixtral_llm,

| | | ```
prompt=dish_prompt_template,
  output_key='recipe'
)
// Third chain - estimates cooking time
recipe_chain = LLMChain(
  llm=mixtral_llm,

  prompt=recipe_prompt_template,
  output_key='time'
)
// Combine into sequential chain
overall_chain = SequentialChain(
  chains=[location_chain, dish_chain, recipe_chain],

  input_variables=['location'],
  output_variables=['meal', 'recipe', 'time'],
  verbose=True
)
``` |
| | | Copied!Wrap Toggled! |
| **RunnablePass through** | A component from langchain_core.runnables that allows function chaining to use the 'assign' method, enabling structured multi-step processing. | ```
from langchain_core.runnables import RunnablePassthrough

// Create each individual chain with the pipe operator
``` |

```python
location_chain_lcel = (

PromptTemplate.from_templ
ate(location_template)
    | mixtral_llm
    | StrOutputParser()
)

dish_chain_lcel = (

PromptTemplate.from_templ
ate(dish_template)
    | mixtral_llm
    | StrOutputParser()
)

time_chain_lcel = (

PromptTemplate.from_templ
ate(time_template)
    | mixtral_llm
    | StrOutputParser()
)

overall_chain_lcel = (

RunnablePassthrough.assig
n(meal=lambda x:
location_chain_lcel.invoke({“
location”: x[“location”]}))
    |
RunnablePassthrough.assig
n(recipe=lambda x:
dish_chain_lcel.invoke({“me
al”: x[“meal”]}))
    |
RunnablePassthrough.assig
n(time=lambda x:
time_chain_lcel.invoke({“reci
```

		pe": x["recipe"]}))
		)
		
		// Run the chain
		result = overall_chain_lcel.invoke({"location": "China"})
		pprint(result)
		
		Copied!Wrap Toggled!
| **Tool** | A class from langchain_core.tools that represents an interface that an agent, chain, or LLM can use to interact with the world. Tools perform specific tasks like calculations and data retrieval. | ```python
from langchain_core.tools import Tool

from langchain_experimental.utilities import PythonREPL


python_repl = PythonREPL()

python_calculator = Tool(
    name="Python Calculator",
    func=python_repl.run,
    description="Useful for when you need to perform calculations or execute Python code. Input should be valid Python code."
)

result = python_calculator.invoke("a = 3; b = 1; print(a+b)")
```<br><br>Copied!Wrap Toggled! |
| **@tool decorator** | A decorator from langchain.tools that simplifies | ```python
from langchain.tools import tool
``` |

| | the creation of custom tools. This tool automatically converts a function into a Tool object. | ```python
@tool
def search_weather(location: str):
    """Search for the current weather in the specified location."""
    # In a real application, this function would call a weather API
    return f"The weather in {location} is currently sunny and 72°F."
```
Copied!Wrap Toggled! |
|---|---|---|
| **create_react_agent** | A function from langchain.agents that creates an agent following the ReAct (Reasoning + Acting) framework. This function takes an LLM, a list of tools, and a prompt template as input and returns an agent that can reason and select tools to accomplish tasks. | ```python
from langchain.agents import create_react_agent


agent = create_react_agent(
    llm=mixtral_llm,
    tools=tools,
    prompt=prompt
)
```
Copied!Wrap Toggled! |
| **AgentExecutor** | A class from langchain.agents that manages the execution flow of an agent. This class handles the orchestration between the agent's reasoning and the actual tool execution. | ```python
from langchain.agents import AgentExecutor


agent_executor = AgentExecutor(
    agent=agent,
    tools=tools,
    verbose=True,
``` |

		handle_parsing_errors=True) result = agent_executor.invoke({"input": "What is the square root of 256?"}) Copied!Wrap Toggled!

Author

Hailey Quach