

Cheat Sheet: Foundations of Generative AI and LangChain

Estimated time needed: 10 minutes

Package/Method	Description	Code Example
pip install	Installs the necessary Python libraries required for the course.	<pre>%%capture !pip install "ibm-watsonx-ai==1.0.8" --user !pip install "langchain==0.2.11" --user !pip install "langchain-ibm==0.1.7" --user !pip install "langchain-core==0.2.43" --user Copied!Wrap Toggled!</pre>
warnings	Suppresses warnings generated by the code to keep the output clean.	<pre>import warnings warnings.filterwarnings('ignore') Copied!Wrap Toggled!</pre>
WatsonxLLM	Facilitates interaction with IBM's Watsonx large language models.	<pre>from langchain_ibm import WatsonxLLM granite_llm = WatsonxLLM(model_id="ibm/granite-3-2-8b-instruct", url="https://us-south.ml.cloud.ibm.com",</pre>

		<pre> project_id="skills-network", params={ "max_new_tokens": 256, "temperature": 0.5, "top_p": 0.2 }) </pre> <p>Copied! Wrap Toggled!</p>
llm_model	Invokes IBM Watsonx LLM with a given prompt and parameters.	<pre> def llm_model(prompt_txt, params=None): model_id = "ibm/granite-3-2-8b- instruct" default_params = { "max_new_tokens": 256, "temperature": 0.5, "top_p": 0.2 } if params: default_params.update(params) granite_llm = WatsonxLLM(model_id=model_id, url="https://us- south.ml.cloud.ibm.com", project_id="skills-network", params=default_params) </pre>

		<pre>response = granite_llm.invoke(prompt_txt) return response</code></pre></pre> <p>1.</p> <p>Copied!Wrap Toggled!</p>
GenParams	A class from the <code>ibm_watsonx_ai.metanames</code> module that provides parameters for controlling text generation, including <code>max_new_tokens</code> , <code>min_new_tokens</code> , <code>temperature</code> , <code>top_p</code> , and <code>top_k</code> .	<pre>from ibm_watsonx_ai.metanames import GenTextParamsMetaNames as GenParams // Get example values GenParams().get_example_values () // Use in parameters parameters = { GenParams.MAX_NEW_TOKENS: 256, GenParams.TEMPERATURE: 0.5, }</pre> <p>Copied!Wrap Toggled!</p>
Basic Prompt	The simplest form of prompting, in which you provide a short text or phrase to the model without special formatting or instructions. The model then generates a continuation based on	<pre>params = { "max_new_tokens": 128, "min_new_tokens": 10, "temperature": 0.5, "top_p": 0.2, "top_k": 1 }</pre>

	<p>patterns it has learned during training.</p>	<pre>prompt = "The wind is" response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n")</pre> <p>Copied!Wrap Toggled!</p>
Zero-shot Prompt	<p>A technique in which the model performs a task without any examples or prior specific training on that task. This approach tests the model's ability to understand instructions and apply its knowledge to a new context without demonstration.</p>	<p>prompt = """Classify the following statement as true or false:</p> <p>'The Eiffel Tower is located in Berlin.'</p> <p>Answer:</p> <p>"""</p> <pre>response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n")</pre> <p>Copied!Wrap Toggled!</p>
One-shot Prompt	<p>Provides the model with a single example of the task before asking it to perform a similar task. This technique gives the model a pattern to follow, improving its understanding of the desired output format and style.</p>	<pre>params = { "max_new_tokens": 20, "temperature": 0.1, }</pre> <p>prompt = """Here is an example of translating a sentence from English to French:</p>

		<p>English: "How is the weather today?"</p> <p>French: "Comment est le temps aujourd'hui?"</p> <p>Now, translate the following sentence from English to French:</p> <p>English: "Where is the nearest supermarket?"</p> <p>"""</p> <pre>response = llm_model(prompt, params)</pre> <p>Copied! Wrap Toggled!</p>
Few-shot Prompt	<p>Extends the one-shot approach by providing multiple examples (typically 2-5) before asking the model to perform the task. These examples establish a clearer pattern and context, helping the model better understand the expected output format, style, and reasoning.</p>	<pre>params = { "max_new_tokens": 10, }</pre> <p>prompt = """Here are few examples of classifying emotions in statements:</p> <p>Statement: 'I just won my first marathon!'</p> <p>Emotion: Joy</p> <p>Statement: 'I can't believe I lost my keys again.'</p> <p>Emotion: Frustration</p> <p>Statement: 'My best friend is moving to another country.'</p> <p>Emotion: Sadness</p> <p>Now, classify the emotion in the following statement:</p>

		<p>Statement: 'That movie was so scary I had to cover my eyes.'</p> <p>“””</p> <pre>response = llm_model(prompt, params)</pre> <p>Copied!Wrap Toggled!</p>
Chain-of-thought (CoT) Prompting	<p>Encourages the model to break down complex problems into step-by-step reasoning before arriving at a final answer. By explicitly showing or requesting intermediate steps, this technique improves the model's problem-solving abilities and reduces errors in tasks requiring multi-step reasoning.</p>	<pre>params = { "max_new_tokens": 512, "temperature": 0.5, }</pre> <p>prompt = “””Consider the problem: 'A store had 22 apples. They sold 15 apples today and got a new delivery of 8 apples. How many apples are there now?'</p> <p>Break down each step of your calculation</p> <p>“””</p> <pre>response = llm_model(prompt, params)</pre> <p>Copied!Wrap Toggled!</p>
Self-consistency	<p>An advanced technique where the model generates multiple independent solutions or answers to the same problem, then evaluates these different approaches to determine</p>	<pre>params = { "max_new_tokens": 512, }</pre> <p>1.</p>

	<p>the most consistent or reliable result. This method helps improve accuracy by leveraging the model's ability to approach problems from different angles.</p>	<p>prompt = """When I was 6, my sister was half of my age. Now I am 70, what age is my sister?</p> <p>Provide three independent calculations and explanations, then determine the most consistent result.</p> <p>"""</p> <pre>response = llm_model(prompt, params)</pre> <p>Copied! Wrap Toggled!</p>
PromptTemplate	<p>A class from <code>langchain_core.prompts</code> module that acts as a reusable structure for generating prompts with dynamic values. It allows you to define a consistent format while leaving placeholders for variables that change with each use case.</p>	<pre>from langchain_core.prompts import PromptTemplate</pre> <p>template = """Tell me a {adjective} joke about {content}.""" prompt = PromptTemplate.from_template(template) // Format the prompt formatted_prompt = prompt.format(adjective="funny", content="chickens")</p> <p>Copied! Wrap Toggled!</p>
RunnableLambda	<p>A class from <code>langchain_core.runners</code> that wraps a Python function into a LangChain</p>	<pre>from langchain_core.runners import RunnableLambda</pre>

	<p>runnable component. It's used to create transformation steps in a chain, especially for formatting or processing data.</p>	<pre>// Define a function to ensure proper formatting def format_prompt(variables): return prompt.format(**variables) // Use in a chain joke_chain = (RunnableLambda(format_prompt) llm StrOutputParser()) Copied!Wrap Toggled!</pre>
StrOutputParser	<p>A class from <code>langchain_core.output_parsers</code> that simply extracts string outputs from LLM responses. It's commonly used as the final step in a LangChain chain to ensure a clean string is returned.</p>	<pre>from langchain_core.output_parsers import StrOutputParser // Create a chain that returns a string chain = (RunnableLambda(format_prompt) llm StrOutputParser()) // Run the chain response = chain.invoke({"variable": "value"}) Copied!Wrap Toggled!</pre>
LCEL Pattern	<p>LangChain Expression Language (LCEL) is a pattern for building</p>	<pre>// Basic LCEL pattern chain = (</pre>

LangChain applications using the pipe operator (|) for more flexible composition. It offers better composability, clearer visualization of data flow, and more flexibility when constructing complex chains.

```
RunnableLambda(format_prompt)  
# Format input
```

```
    | llm          # Process with  
    LLM  
    | StrOutputParser()  # Parse  
    output  
)
```

```
// Run the chain  
result = chain.invoke({“variable”:  
“value”})
```

```
// More complex example  
template = “””
```

Answer the {question} based on the {content}.

Respond “Unsure about answer” if not sure.

Answer:

```
“””  
prompt =  
PromptTemplate.from_template(t  
emplate)  
qa_chain = (
```

```
RunnableLambda(format_prompt)  
| llm  
| StrOutputParser()  
)
```

```
answer = qa_chain.invoke({  
“question”: “Which planets are
```

rocky?",
 "content": "The inner planets
are rocky."
})

Copied! Wrap Toggled!

Author

Hailey Quach

